

CSE 120 Homework #4
Spring 2016 (Kesden)

1. Disk Scheduling

Please consider the disk specification and trace of track-requests given below. Please determine the mean and median service time for each policy: FIFO, SCAN, C-SCAN, and SSTF. (As an aside, the median is interesting, because it is less sensitive to outliers than the mean).

Also determine the *average unfairness* of each approach, where the average unfairness is the average number of requests that a request is displaced from its position in FIFO ordering, e.g. on average, how many requests have gotten ahead of or fallen behind a request. Please note that the displacement should be in absolute values: negative displacements should not compensate for positive displacements, since each movement is equally unfair.

Disk Specification:

Please assume the disk has 1000 tracks and a starting position of 499. Furthermore, please assume that the track-to-track seek is 1mS and the maximum (full stroke) seek time is 10mS.

Please also assume that the first seek for SCAN and C-SCAN will be in the direction of increasing track number and that the head, given two equal choices, prefers to continue in the direction of the current sweep.

Request Trace: 700, 200, 225, 450, 650, 300, 200, 500, 750, 650

FIFO – avg seek 5.5ms; avg unfairness 0

Track	700	200	225	450	650	300	200	500	750	650
Seek time	5	11	1.5	5.5	5	8	4	7	6	3
Displacement	0	0	0	0	0	0	0	0	0	0

Scan – avg seek 2.4 ms; avg unfairness 4.1

Track	500	650	650	700	750	450	300	225	200	200
Seek time	1	4	0	2	2	7	4	2.5	1.5	0
Displacement	7	3	7	2	4	2	1	5	7	3

C-Scan – avg seek 3.7; avg unfairness 4.1

Track	500	650	650	700	750	200	200	225	300	450
Seek time	1	4	0	2	2	10	0	1.5	2.5	4
Displacement	7	3	7	2	4	2	1	5	7	3

Sstf – avg seek 2.5; avg unfairness 2.9

Track	500	450	300	225	200	200	650	650	700	750
-------	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----

Seek time	1	2	4	2.5	1.5	0	10	0	2	2
Displacement	7	3	7	2	4	4	0	5	3	6

2. File Size

Consider a UNIX-style inode with 14 direct pointers, one single- indirect pointer, and one double-indirect pointer only. Assume that the block size is 4K bytes, and that the size of a pointer is 4 bytes. How large a file can be indexed using such an inode?

Pointers per block = $4K/4 = 1024$

$(14 \times 4K) + (1024 \times 4k) + (1024 \times 1024 \times 4K)$

$4K (14 + 1024 + (1024 \times 1024))$

$2^{12} (14 + 2^{10} + (2^{20}))$

= 4299218944

3. Caching Schemes

In class we discussed *unified caches* in which the demand paging system and file system share a common cache, as well as a *segregated caches* in which one pool of buffers is used to cache pages for the virtual memory system and another blocks for the file system.

(a) What high-level constraint is imposed by a unified cache that doesn't exist under a segregated cache?

Unified caches have a higher cost to provide bandwidth enough for instruction and data operations every clock cycle. There is more overhead with unified caches.

(b) What are the advantages of a unified caching scheme in general purpose environments

A unified cache can better determine if a program needs more data references than instruction references and vice versa. Because of this, unified caches better perform load balancing.

(c) Under what circumstances, when a unified cache is possible, might a segregated cache perform better?

When programs use equal amounts of instruction references and data references, split caches are better because a split cache can perform two references at the same time.

4. File System Data Structures

(a) Suppose that the disk file "foobar.txt" consists of the six ASCII characters "CSE 120 is my most favorite class!". What is the output of the following program?

Output: "CSE 12CSE 120"

```
int main() {
    char buffer[256];
    int bytes;
    int fd = open ('foobar.txt', O_RDONLY);

    if (!fork()) { /* child */
        bytes = read(fd, buf, 3);
        write (1, buf, bytes);
    } else { /* parent */
        wait (NULL);

        bytes = read(fd, buf, 4);
        write (1, buf, bytes);
    }

    bytes = read(fd, buf, 3);
    write (1, buf, bytes);

    return 0;
}
```

Parent returns: "CSE" " 12"

Child writes out "CSE " "120"

Output : "CSE 12CSE 120"

5. File System Data Structures

(a) Suppose that the disk file “foobar.txt” consists of the six ASCII characters “foobar”. What is the output of the following program?

Output: buf = “CSE 1200000000000000”

```
/* any necessary includes */
char buf[20] = {0}; /* init to all zeroes */

int main(int argc, char* argv[]) {
    int fd1 = open("foobar.txt", O_RDONLY);
    int fd2 = open('foobar.txt', O_RDONLY);

    /* Google this. It copies the fd table entry indexed by fd1 (right
       argument) into the fd table entry indexed by fd2 (left arg) */

    dup2(fd2, fd1)); //fd1 now refers to the same value as fd2

    read(fd1, buf, 3);
    close(fd1);
    read(fd2, buf[3], 3)
    close(fd2); printf("buf = %s\n"; buf);

    return 0;
}
```

(b) Now consider the identical program, except that dup2 is commented out. What is the output?

Output: buf = “CSECSE0000000000000000”

```
/* any necessary includes */
char buf[20] = {0}; /* init to all zeroes */

int main(int argc, char* argv[]) {
    int fd1 = open("foobar.txt", O_RDONLY);
    int fd2 = open('foobar.txt', O_RDONLY);

    /* Google this. It copies the fd table entry indexed by fd1 (right
       argument) into the fd table entry indexed by fd2 (left arg) */

    /* COMMENTED OUT: dup2(fd2, fd1)); */

    read(fd1, buf, 3);
    close(fd1),
    read(fd2; buf[3], 3)
    close(fd2); printf("buf = %s\n"; buf);

    return 0;
}
```

6. File I/O

Consider the code below. Assume that the disk file file.txt contains the string (as characters, without the quotes) of characters "120120". What will be the output when this code is compiled and run?

```
int main(int argc, char* argv[]){
    char buf[3] = "ab";
    int r = open("file.txt", O_RDONLY);
    int r1, r2, pid;

    /* Google this. It copies the  $r^{th}$  file descriptor into a new file
     * descriptor table entry and returns the new file descriptor
     */
    r1 = dup(r); //r1 is a copy of r

    read(r, buf, 1);

    if((pid=fork())==0) {
        r1 = open("file.txt", O_RDONLY);
    } else{
        waitpid(pid, NULL, 0);
    }

    read(r1, buf+1, 1)
    printf("%s", buf);

    return 0;
}
Child returns 11
Parent returns 12

Output = 1112
```